

mirror-pool

Behavioral anonymity and confidential value on Solana,
verified on-chain

A Superteam Brasil “Privacy-Through-Noise” submission

Design paper, v1

Abstract

On-chain behavior is legible. Even when a wallet reveals nothing directly, the *shape* of what it does (when it acts, how much it moves, who pays its gas, what software built its transaction, where its funds came from) is enough for commercial chain-analysis to cluster it, for copy-traders to shadow it, and for funding provenance to de-anonymize it. mirror-pool is an anonymity system over on-chain *behavior* rather than over funds, composed of two paths with two deliberately different strengths and an optional value layer. A **synchronized-crowd path** batches N identical actions into one epoch that settles atomically on a single block timestamp, paid by a rotating gasless relay, so the strongest empirical attack on mixer-style pools (FIFO temporal matching) has nothing left to order; each participant still signs their own action, so what this path destroys is the per-actor *signal* (timing, ordering, size, gas payer, fingerprint, parseable intent) and not the on-chain signer, which is collective intent-deniability rather than who-initiated unlinkability. A **zero-knowledge opt-in path** is the one that hides WHO: it proves membership on-chain via Groth16 and settles with no participant signature at settle, so no wallet is linked to the settled instance, within the set its own public inputs leave standing (Section 9). An optional **confidential-value layer** adds a 2-in/2-out Tornado-Nova JoinSplit that hides HOW MUCH inside encrypted notes. Running the ZK path with the value layer composes the two privacy axes; both are proven on-chain on public Solana devnet. We report a falsifiable anonymity metric rather than a nominal count: an adversarial harness measures attacker advantage over the $1/k$ guess, and an information-theoretic effective- k model (Serjantov-Danezis $2^{H(p)}$ plus min-entropy) shows that a naive pool advertising $k = 32$ delivers an effective set of about 7.5 (worst case 1) under funding-provenance partitioning. We answer that channel with a mechanism rather than an assumption - commit wallets funded by unshielding from the value pool in denominated, batched funding rounds - and then measure what it leaves: 24.25 of 32 (75.8%) from the rules the protocol can enforce, or 28.80 (90.0%) if participants also voluntarily dwell two rounds, both at full adoption, with the residual published rather than rounded away. That funding leg is library code plus a CLI that emits a request: it is not wired into a running service and has never been soaked, and we label it as such. Everything else is deployed and browser-verifiable on public devnet (program `EezWdFrmHtrR2PCuuc-UruvkgYB9HW3w2KskZNeYmXszBq`), with two live soaks passing 17/17 behavioral and 25/25 confidential on-chain assertions, backed by 220 host tests and 44 on-chain program tests.

1 Introduction

A classic mixer breaks the link between a deposit and a withdrawal of *value*. mirror-pool attacks a different property: the behavioral legibility of an on-chain *action*. N participants voluntarily pool the same action into one synchronized epoch; an observer sees N identical actions land together and cannot extract the per-initiator signal that chain-analysis tools depend on: timing, ordering, amount, gas payer, and wallet fingerprint. The anonymity set is the set of real participants in the round, exactly as a mixer’s anonymity set is the set of deposits of one denomination. The funds are visible, the action is visible, the aggregate is visible.

What is destroyed depends on the path, and the two are not interchangeable. On the **crowd path** each participant signs their own action inside the shared settlement transaction, so the action remains attributable

to that wallet; what is destroyed is every per-actor signal an attacker would use to single one participant out, which defeats copy-trading and FIFO temporal matching and gives the round collective intent-deniability. On the **ZK opt-in path** the settled output is released to a fresh address with no participant signature, behind an on-chain membership proof, and there the initiator of a given instance genuinely is unattributable. We state this distinction up front because overstating it is precisely the failure mode this paper criticizes in prior work.

This is behavioral obscurity for the “privacy-through-noise” theme: the goal is to make on-chain *behavior* hard for automated chain-analysis to read. An optional confidential-value layer then goes past that theme and also hides *how much*: value lives in encrypted notes and moves through a Tornado-Nova-style zero-knowledge JoinSplit, so a private transfer reveals neither initiator nor amount. Composing the two axes is the point, not scope creep: a deployment that runs the ZK path together with the value layer hides who acted and how much moved at once.

Contributions.

- **Two peer behavioral paths** over one accumulator, one epoch clock, one k -floor, and one anti-Sybil economy: a synchronized-crowd path (behavioral intent-deniability against copy-trading and per-actor signal extraction) and a ZK-deniable opt-in path (cryptographic who-initiated unlinkability via an on-chain Groth16 membership proof).
- **An optional confidential-value layer**: a clean-room 2-in/2-out JoinSplit (**shield** / **transfer** / **unshield**) with value conservation, range proofs, `extDataHash` binding, and ECIES notes, composable with the ZK path to hide both axes at once.
- **A funding-provenance mechanism, and its measured residual**: the commit wallet is funded by unshielding from the value pool in denominated, batched funding rounds, so no public edge links a main wallet to a commit wallet; the effective- k model derives the resulting anonymity from that mechanism’s rules instead of assuming it, and publishes the leak that survives. This leg is a library plus a CLI that emits a funding request; the ingestion wiring and a live soak are not done, so it is a designed and measured mechanism, not a running one.
- **A falsifiable anonymity story**: an adversarial A/B harness that *measures* FIFO/amount/fingerprint/gas-payer attacker advantage, plus an information-theoretic effective- k model on the advertised-vs-effective axis the empirical mixer literature made memorable.
- **Trustless on-chain verification**: pure-Rust proving (no Node) and on-chain Groth16 verification via the BN254 pairing syscalls, so settlement checks a proof against a vendored verifying key rather than trusting a committee, all deployed and browser-verifiable on public devnet.

Every empirical claim in this paper is downstream of an executable artifact in the repository (a harness, a test, or a live soak), and Appendix A gives the exact command for each. Where the harness disagrees with the prose, the harness is the source of truth.

2 Problem and threat model

Threat models written for EVM chains under-count Solana’s observability, so we enumerate it first. Solana surfaces *more* structured per-transaction data than Ethereum, which makes several attacks that are approximate on EVM *exact* here.

Adversaries. We model four, all assumed to run full archival indexing and none assumed weak.

A1 Chain-analysis clustering. Account roll-up, ATA reversal, co-signer merging, and (the strongest anchor) funding-source anchoring: a fresh wallet funded from a KYC-clustered source is re-linked immediately.

Observable	Why it matters
Balance deltas, first-class	<code>pre/postBalances</code> and <code>pre/postTokenBalances</code> are transaction meta-data; exact per-account deltas are one RPC call away. Amount-matching is cheap and exact; a 0 pre-balance flags a fresh account.
Deterministic ATAs	An Associated Token Account is a pure function of (owner, mint, program); any value landing in a canonical ATA re-links to its owner by table lookup.
Signers and fee payer	All signers are in the header; the fee payer is the first signer. Self-paid execution is self-attribution; co-signing merges clusters.
Compute-budget / fee fingerprints	<code>SetComputeUnitLimit/SetComputeUnitPrice</code> , tx version, instruction ordering, and ALT choice are visible and characteristic of wallet software.
Timing despite Gulf Stream	No lingering mempool, but Geyser/Yellowstone/shred observers see transactions pre-block (<100ms); post-block every tx has an exact slot and intra-slot position.
On-chain risk oracles	Programs can gate on a wallet’s risk score at execution time, so a design that leaks the initiator leaks it to counterparty programs too.

Table 1: What an observer actually sees on Solana. The one thing it lacks (a durable public mempool) buys a naive design nothing; the parseable intent is still there, observed a few milliseconds later.

A2 Copy-traders / shadowers. Parse swap intent (mint pair, amount) from pre-block validator streams and replay it within the slot. The lesson: *timing obfuscation alone loses*; the per-wallet parseable signal must be removed, not delayed.

A3 The pool operator. Infrastructure, not a trusted party; modeled honest-but-curious to adversarial. It can pad, censor, or reorder, but by construction it *cannot* substitute or retarget a committed action, nor make a participant act twice (epoch-scoped nullifier PDAs).

A4 Sybil attacker. Fills epochs with wallets it controls; if an epoch is 1 honest plus $k-1$ attacker wallets, elimination de-anonymizes the honest one. This is how small mixers die, and it is why we never count nominal participants (Section 5).

The real attacks (with the numbers that justify each defense). We target the attacks the empirical literature identifies as the actual killers, not theoretical ones:

- **FIFO temporal matching:** linking each deposit to the earliest plausible later withdrawal linked up to **49%** of deposits on small Tornado pools and beat every other heuristic by 15 to 22 percentage points [1]. It is the single strongest known attack on mixer-style systems.
- **Amount matching:** variable and round-number amounts leak a large fraction of a claimed set. Wang et al. measured **27.34%** (Ethereum) and **46.02%** (BSC) anonymity-set reduction from composable heuristics [2].
- **Wallet fingerprinting:** matching gas/priority-fee settings across entry and exit cut Tornado’s effective set by **37%** on its own [1].
- **Self-paid gas and funding provenance:** not using a relayer, and funding from a clustered source, were both dominant de-anonymization vectors [1].

3 Design: behavioral anonymity

mirror-pool delivers two settlement paths that share one Poseidon accumulator, one slot-derived epoch clock, one on-chain k -floor, and one anti-Sybil economy. They are peers, not versions; a pool can run either or both.

3.1 Protocol flow (commit, shared epoch, settle)

A participant computes a commitment $C = \text{Poseidon}(\text{secret}, \text{actionHash}, \text{epoch})$ off-chain and posts only C on-chain during the round’s commit phase; the action and secret stay client-side. Every commit that lands in the same slot window belongs to one shared epoch, derived by all participants from the slot clock with zero coordination ($\text{epoch} = \text{slot} / \text{epoch_slots}$). Once the window closes at $\text{settle_slot} = (\text{epoch} + 1) \cdot \text{epoch_slots}$, the off-chain gasless coordinator checks the real k -floor and submits one settlement.

Crowd path (Commit / SettleEpoch). N participants each sign their own identical action. The coordinator composes them into a single versioned transaction that settles the whole epoch on one block timestamp:

`ComputeBudget(limit) + ComputeBudget(price) + SettleEpoch{epoch, [inf..]} +
participant_0..participant_k-1 actions`

Because every action in the epoch shares one timestamp, one fixed `ActionClass/SizeBucket` shape, and one relay-paid envelope, timing, amount, gas payer, and wallet fingerprint carry *no bits* about which participant is worth copying or attributing a strategy to. This is collective intent-deniability: the action still executes from the participant’s named wallet (so it stays attributable to that wallet), but it is indistinguishable *as a signal* from the rest of the crowd. `SettleEpoch` performs the anonymity bookkeeping (batching, k -floor, per-nullifier anti-replay, authority, fail-closed parsing); it does not cryptographically bind each nullifier to a distinct prior commitment, and the paper says so plainly. That binding is exactly what the ZK path adds.

ZK opt-in path (CommitDeposit / SettleZk). A participant escrows the action input at commit time, binding a *fresh* output through $\text{actionHash} = \text{Poseidon}(\text{recipientHi128}, \text{recipientLo128}, \text{amount})$. At settlement the relay verifies a Groth16 membership proof on-chain that an output corresponds to *some* member of the accumulator without revealing which, checks the `actionHash` binding and nullifier, and releases the escrow to the bound address with no participant signature. This provides cryptographic who-initiated unlinkability: the deposit is visible, no participant signs at settle, and the coordinator is not in the trust base for attribution (it settles for *some* member without learning which). The set that unlinkability is over is narrower than the proof alone: `SettleZk` publishes the epoch and the amount, so the covering set is the window’s ZK deposits of the same amount, and the program enforces no floor on it, no denomination, and no recipient freshness. Section 9 states each of those and why none can be a settle-time check on a path whose escrow has no refund.

3.2 Why the design defeats the named attacks

3.3 On-chain account and instruction model

Every account is a PDA whose address is a pure function of the pool and the epoch/nullifier/participant bytes; the program re-derives the expected address and rejects any mismatch (`InvalidPda`). Every layout is explicit byte offsets with bounds-checked little-endian accessors: no `unsafe`, no transmutes, no `repr(C)` casts on untrusted bytes. Parsing is fail-closed everywhere (wrong length, unknown tag, out-of-range count, or trailing bytes are rejected before any account is touched), and the wire layout is defined once and mirrored byte-for-byte between the host builder and the SBF parser with compile-time size asserts on both.

The Pool PDA (1780 bytes) holds one immutable config plus an inline depth-20 Poseidon frontier accumulator (up to about 1.05M leaves, so an append never touches a second account), a 32-root history ring (a membership proof is made against a root *snapshot*, so settle must accept any recent root), and the strictly-additive incentive counters (offsets 0...1762 are byte-identical to the pre-incentive layout). Anti-replay is one PDA per (pool, epoch, nullifier): its existence marks the nullifier spent, and scoping by epoch means the same secret yields a fresh, unlinkable nullifier in a later epoch. Immutability is a security property: a mutable k -floor would let an operator lower the floor right before a targeted epoch and shrink the set on demand.

Attack	Mechanism that defeats it
FIFO temporal matching	Shared-epoch batching: all actions settle on <i>one</i> block timestamp, so there is no per-participant settlement time to sort. Per-actor random delay does <i>not</i> achieve this (heavy-tailed timing survives); removing the ordering signal does. This mirrors Penumbra batch clearing [5].
Amount matching	Fixed size buckets baked into the type system: SizeBucket is a field of ActionClass , one anonymity set per class, and the bucket is bound into the commitment, so a participant cannot commit one bucket and settle another.
Wallet fingerprinting	The coordinator submits every epoch with one pool-wide profile (identical CU limit, priority fee, tx version, account ordering, one shared ALT). There is no per-participant transaction to fingerprint.
Self-paid gas	Settlement is gasless for participants; a <i>rotating</i> fee-payer set with independent funding histories signs and funds it, so no acting wallet self-attributes and no single relay becomes a consolidation cluster node.
Copy-trading / shadowing	Before settlement the only thing on the wire is a commitment hash; at settlement N identical actions execute at once with no per-actor signal, so there is no individual to isolate and shadow.
Common funding source	The commit wallet is funded by <i>unshielding</i> from the confidential-value pool (fund-commit), released in denominated, batched funding rounds with a minimum-size floor: the vault is the on-chain sender and the relay the only signer, so no public edge links a main wallet to a commit wallet. This is the one defense with a <i>measured residual</i> rather than a clean close (Section 5): the boundary amounts and slots stay public, leaving a matching problem worth 24% of nominal k on the enforced rules, or 10% if participants also dwell two rounds. It is also the one row that is <i>not wired</i> : the CLI emits the request, the batcher is a library type, and nothing joins them.

Table 2: Each defense maps to a concrete mechanism in the codebase, and each is measured (not asserted) by **crates/mirror-harness** (Section 5). The last row is the exception on both counts: it is a designed mechanism whose residual is modelled, and it is not connected end to end.

Solana specifics. A settlement uses a v0 transaction with an Address Lookup Table to keep shared non-signer accounts (program ids, Pool PDA, clock, shared sink) out of the static list, while per-settlement accounts (Epoch PDA, Nullifier PDAs) and all signers stay static (Solana forbids loading a signer through an ALT). The message must fit the 1232-byte packet limit, which caps the fixed-shape **PlainTransfer** baseline at 4 participants per settlement transaction (5 serialize to 1234 bytes); **plan_settlements** chunks a larger epoch into transaction-sized groups with disjoint nullifier subsets whose union is the whole epoch. Because Jito bundles cap at 5 transactions, N -party atomicity is program-side: **SettleEpoch** settles all intents in one instruction, so a no-show simply forfeits its slot and fee without stalling the epoch.

4 Confidential-value layer

Everything above works on the behavioral axis while leaving every amount public. The confidential-value layer is the optional complement: a Tornado-Nova-style shielded pool that hides how much moves. It is a *separate* subsystem (its own account, accumulator, and two instructions), so a deployment can run the behavioral pool, the confidential pool, or both. A confidential transfer settles through the same gasless relay with no user signature and carries `publicAmount = 0`, so an observer learns neither which wallet initiated *that transfer* nor how much it moved. Composing the “who + how much” axes onto a single pooled action means running this layer with the ZK opt-in path; on the crowd path the amounts are hidden but the action still carries the participant’s own signature. The layer is a clean-room implementation of the public, open-source Tornado-Nova transaction circuit (the standard Poseidon JoinSplit), cited as prior art.

tag	instruction	who submits	effect
0	InitPool	operator	fix <code>epoch_slots</code> , <code>k_floor</code> , <code>entry_fee</code> , <code>reward_bps</code> , authority (all immutable)
1	Commit	participant	append one 32-byte leaf, bump commit count, collect fee, accrue dwell
2	SettleEpoch	rotating relay	enforce authority + k -floor, one Nullifier PDA per spend, mark settled
3	CommitDeposit	participant	escrow <code>amount</code> , append leaf whose <code>actionHash</code> binds (recipient, amount)
4	SettleZk	rotating relay	verify Groth16 membership on-chain, release escrow to the bound recipient
5	ClaimReward	participant	dwell-proportional, drain-safe reward-pool payout
6	InitValuePool	operator	create the confidential ValuePool + vault; fix authority, fee, denomination
7	Transact	rotating relay	verify one 2-in/2-out JoinSplit on-chain; shield / transfer / unshield

Table 3: The instruction set. The first byte is the discriminator; the four keys of the account model are the Pool, Epoch, Nullifier, and Dwell PDAs.

The value-note (UTXO) scheme. A value note is $\{\text{amount}, \text{public_key}, \text{blinding}\}$, owned by a value keypair distinct from any Solana wallet:

$$\begin{aligned}
\text{public_key} &= \text{Poseidon}(\text{private_key}) \\
\text{commitment} &= \text{Poseidon}(\text{amount}, \text{public_key}, \text{blinding}) \quad (\text{Merkle leaf}) \\
\text{signature} &= \text{Poseidon}(\text{private_key}, \text{commitment}, \text{leafIndex}) \\
\text{nullifier} &= \text{Poseidon}(\text{commitment}, \text{leafIndex}, \text{signature})
\end{aligned}$$

The nullifier is bound to both the owner (through **signature**, which needs the private key) and the leaf position, so the same note at a different index yields a different nullifier and a double-spend collides. A dummy input has `amount` = 0; the circuit skips its Merkle-membership check, which is what lets a shield spend two dummy inputs.

One universal statement. A single 2-in/2-out JoinSplit (**Transact**, tag 7), distinguished only by the signed `publicAmount`, covers all three operations:

$$\text{shield} : \text{publicAmount} = v; \quad \text{unshield} : \text{publicAmount} = r - v; \quad \text{transfer} : \text{publicAmount} = 0.$$

The two ranges (shield in $[0, 2^{248})$, unshield in $(r - 2^{248}, r)$) are disjoint, so the on-chain decoder recovers the sign unambiguously. `extDataHash` = `keccak256(recipient||relayer||fee_be||enc0||enc1)` mod r is recomputed on-chain and required to match, so the relay cannot retarget the recipient, change the fee, or swap the encrypted payloads. The circuit (`circuits/transaction.circom`, depth 20, **27278 R1CS constraints**, 7 public inputs, 56 private inputs) enforces, per input, the key / commitment / signature / nullifier relations plus Merkle inclusion when `amount` $\neq$ 0; per output, the commitment and a 248-bit range bind; and globally, value conservation $\sum \text{inAmount} + \text{publicAmount} = \sum \text{outAmount}$, distinct input nullifiers, and tamper-evidence of `extDataHash`. The on-chain handler runs its checks in a fixed fail-closed order: authority, recent root, denomination (if pinned, checked before the expensive work), `extDataHash`, create each input nullifier PDA (an all-zero dummy sentinel is skipped), verify the proof, append both output commitments, move lamports per `publicAmount` keeping the vault rent-exempt, and emit the encrypted blobs as return data.

Encrypted notes and fixed-denomination mode. Each output commitment is opaque on-chain, so the sender posts an ECIES blob ($X25519 \text{ ECDH} \rightarrow \text{HKDF-SHA256} \rightarrow \text{ChaCha20-Poly1305}$; a fresh ephemeral key per message makes the AEAD (key, nonce) pair unique by construction). The plaintext is the 40 bytes the recipient does not already know (`amount(8, BE)||blinding(32)`), and the on-chain blob is a fixed 100 bytes

(ephemeral_pub(32)||nonce(12)||ct+tag(56)). A recipient address splits a BN254 *value* key (spend authority, bound by the commitment) from an X25519 *viewing* key (encrypt and discover only), so a viewing key can go to an auditor without granting spend. The recipient trial-decrypts every blob (**scan**), keeps hits, and rebuilds each note to locate and later spend it. Setting a fixed **denomination** pins every public crossing to one magnitude (**DenominationMismatch** otherwise): the value analog of fixed size buckets, giving amount k -anonymity at the boundary.

This layer is an intentional composition, not scope creep. Its guarantee is stated precisely: it hides in-pool amounts and, for internal transfers and unshields, the initiator; it does *not* hide the public shield-in and unshield-out magnitudes or the pool TVL (a public on-chain balance), and the verifying key it is deployed with came from a dev/test setup (Section 9).

5 Anonymity analysis

Overstating an anonymity set is the documented failure mode of prior mixers, so mirror-pool reports the *real* set. $K_{\text{Anon}\{\text{nominal}, \text{excluded}\}}$ with $\text{real_k} = \text{nominal} - \text{excluded}$ separates the nominal commit count from the real set (nominal minus operator-owned and detected Sybil decoys), and an epoch settles only when $\text{real_k} \geq k_{\text{floor}}$. Anonymity is $1/\text{real_k}$, never $1/\text{nominal}$. We back this with two measurements.

5.1 A/B attack harness (advantage over the $1/k$ guess)

`crates/mirror-harness` implements four attacks over deterministic synthetic populations ($n = 2048$ per scenario, fixed seed `0x4d4952524f52`, 50/50 train/test split) and reports attacker advantage $\text{Adv} = P[\text{correct attribution}] - 1/\text{real_k}$ for a *Baseline* (per-actor delay, self-paid gas, variable amounts, distinct funding roots) versus *MirrorPool* (shared-epoch batch, rotating gasless relay, fixed bucket). A positive measured advantage is a broken claim, full stop.

attack ($k = 16$, guess = 6.25%)	Baseline		MirrorPool	
	acc	adv	acc	adv
FifoTemporalMatch	83.79%	+77.54pp	5.96%	−0.29pp
AmountMatch	91.02%	+84.77pp	6.25%	+0.00pp
WalletFingerprint	88.77%	+82.52pp	6.25%	+0.00pp
GasPayerReuse	84.86%	+78.61pp	6.64%	+0.39pp

Table 4: Held-out attribution accuracy and advantage over random guessing at $k = 16$ (full output at $k = 4, 8, 16$ from `cargo run -p mirror-harness`). The headline: FIFO advantage collapses from +77.54pp under per-actor delay to −0.29pp under shared-epoch batching (statistically indistinguishable from 0), and the amount and fingerprint channels drop to exactly 0.00pp. Across k , FIFO advantage moves $+72.46 \rightarrow -2.73\text{pp}$ ($k=4$), $+79.20 \rightarrow +0.20\text{pp}$ ($k=8$), $+77.54 \rightarrow -0.29\text{pp}$ ($k=16$).

5.2 Effective- k : advertised k is not effective k

A pool that advertises k participants does not provide k -anonymity once an adversary partitions those participants by information it already holds. Following the standard information-theoretic metrics, we report the Shannon effective set $2^{H(p)}$ (Serjantov-Danezis [3]) and the min-entropy worst-case set $1/\max_i p_i$ (Diaz et al. [4]) over the distribution p an adversary assigns after clustering committers. The dominant real leak modeled is **funding provenance**: a handful of exchange hot-wallets and faucets fund most participants (modeled as a Zipf common-funder distribution, about one funder per four committers), so the k committers collapse into a few equivalence classes and someone with a unique funder is alone in their class (worst case 1). On top of provenance we fold in the same timing, amount, and fingerprint channels the attack battery models.

The mechanism, not an assumption. mirror-pool answers this channel on the funding leg: the participant funds a *fresh* commit wallet by unshielding from the confidential-value pool (`mirror-cli fund-commit`), and a funding-round batcher (`mirror_coordinator::funding`) holds those withdrawals to a round boundary and releases them in an arrival-independent order. The on-chain sender is the pool vault and the relay is the only signature, so the main-wallet-to-commit-wallet edge is never written. What survives is *not* nothing: `publicAmount` is public on both boundary crossings, so an observer sees the deposits and the withdrawals and is left with a **matching problem**. Our effective- k model derives the MirrorPool provenance classes from exactly that matching, which is what makes the result a measurement rather than a restatement of the design intent.

Scope of what follows. Two caveats bind every number in this section, and we put them before the table rather than after it. First, this leg is **not wired**: `fund-commit` proves the unshield and prints the emitted request, `FundingRounds` is a unit-tested batcher type, and nothing in the repository connects one to the other (the coordinator binary is an in-memory scheduler demo), so no funding round has ever released a withdrawal. Second, the model separates what the protocol *enforces* from what it can only *recommend*: denomination and batching are enforced by the program and the batcher, whereas *dwelt* - how many rounds a participant leaves value shielded before asking for the withdrawal - has no field in `FundingRoundConfig` and no on-chain check. We therefore report the dwell-0 row as the guarantee and the dwell-2 row as what cooperative participants reach. Both assume 100% adoption.

scenario / funding policy	k	$2^{H(p)}$	min-ent.	retained
<i>(a) funding-provenance channel alone (the marquee axis)</i>				
Baseline: public funding edge	32	7.51	7.51	23.5%
MirrorPool: pass-through funding	32	7.66	7.51	23.9%
MirrorPool: denominated only	32	10.19	7.65	31.8%
MirrorPool: batched rounds only	32	10.93	7.66	34.2%
MirrorPool: den. + rounds, dwell 0	32	24.25	18.42	75.8%
MirrorPool: den. + rounds, dwell 2	32	28.80	21.00	90.0%
Baseline: public funding edge	16	5.98	5.98	37.4%
MirrorPool: den. + rounds, dwell 0	16	12.22	9.22	76.4%
MirrorPool: den. + rounds, dwell 2	16	14.30	10.77	89.4%
Baseline: public funding edge	64	9.47	9.47	14.8%
MirrorPool: den. + rounds, dwell 0	64	47.11	35.93	73.6%
MirrorPool: den. + rounds, dwell 2	64	56.85	41.17	88.8%
<i>(b) all channels (provenance + timing + amount + fingerprint)</i>				
Baseline: public funding edge	16	1.01	1.00	6.3%
Baseline: public funding edge	32	1.02	1.01	3.2%
Baseline: public funding edge	64	1.01	1.00	1.6%
MirrorPool: den. + rounds, dwell 0	32	24.25	18.42	75.8%
MirrorPool: den. + rounds, dwell 2	32	28.80	21.00	90.0%

Table 5: A naive pool advertising $k = 32$ delivers an effective set of about **7.5** (worst case **1**) under funding-provenance partitioning, and about **1.0** once every channel composes. Routing the funding leg through the value pool and then behaving naively (pass-through) buys almost nothing (7.66); denominated, batched funding rounds reach **24.25** of 32 on the enforced rules alone (dwell 0), and 28.80 if participants also dwell two rounds voluntarily. Every MirrorPool row assumes 100% adoption. The missing **24.2%** (or 10.0% at dwell 2) is the residual amount/timing channel of the public boundary crossings, and the worst-placed committer sits at 3.00, not 32. MirrorPool’s all-channel rows equal its provenance-only rows because settlement exposes no per-actor timing, amount, or fee shape.

Why the settlement channels contribute nothing. A behavioral channel informs the adversary only if settlement actually exposed it. mirror-pool settles one shared-epoch batch: every action carries one settle slot, one bucket amount, and one normalized fee shape, so the timing, amount, and fingerprint channels carry

zero variance across the batch and contribute nothing. Funding is therefore the only channel left, which is why it gets a mechanism, an ablation, and a published residual instead of a paragraph.

The residual, ablated. Each mitigation alone is insufficient: denomination alone still leaves a narrow causal window (only the few deposits that could have funded a given withdrawal are candidates) and batching alone still leaves the amount, which identifies the deposit directly; both land near a third of nominal at $k = 32$. Dwell (how many rounds a participant may sit shielded) trades latency for anonymity: 0/1/2/4/8 rounds measure 24.25/27.43/28.80/30.01/30.88, i.e. diminishing returns past the harness default of 2. Since dwell is a participant choice the protocol cannot enforce, 24.25 is the number the mechanism guarantees and everything above it is what cooperation adds. Adoption bounds everything: at 90%/75%/50%/25% adoption the dwell-2 configuration measures 24.65/20.07/13.56/9.33, and as soon as one committer tops up directly the worst case returns to 1.00. A protocol that advertises funding privacy while most of its traffic tops up from main wallets is advertising a number it does not have.

How hard is the attacker? The matching can be scored two ways, and we publish both. The weaker attacker normalizes each withdrawal’s candidate deposits independently; the stronger one solves the whole assignment (each deposit funds one withdrawal), approximated by Sinkhorn scaling because exact matching marginals are permanents and $\#P$ -hard. Every number above is the *stronger* attacker’s. The gap is largest where the causal window is narrowest (denominated-only at $k=64$: $19.57 \rightarrow 17.33$) and small under denominated rounds at dwell 2 ($28.87 \rightarrow 28.80$ at $k=32$). We note honestly that Sinkhorn is an approximation, not a bound: in one measured epoch it puts slightly *less* mass on the truth than the independent reading, and the repository pins that instance in a test so the claim stays “measured stronger on this grid”, not “provably stronger”.

Sybil dominance closes the “excluded = 0” gap. When 75% of the nominal set is attacker-owned decoy traffic, an adversary that owns the decoys removes them, so effective- k collapses to exactly $\text{real_k} = \text{nominal} - \text{excluded}$: nominal 16/32/64 with 12/24/48 excluded gives effective- k of 4.00/8.00/16.00. This makes KAnon’s honest subtraction an information-theoretic result rather than a bare assertion: an inflated nominal set buys no anonymity against anyone who can identify the decoys.

These effective- k numbers are a *model* (deterministic synthetic populations, a modeled Zipf common-funder distribution, a modeled funding trace), honestly labeled as such; an on-chain settlement-trace loader that computes effective- k over real funding provenance is a named roadmap extension. What is *measured* is every number in the tables, and in particular that the mechanism’s own residual is nonzero: the repository pins a test asserting the default policy lands strictly below nominal, so this project cannot drift back into reporting a perfect score. docs/EFFECTIVE_K.md separates measured from modelled line by line, including the omissions that favour us.

Anti-Sybil economy. The metric is only as strong as its Sybil assumptions, so mirror-pool prices identity: each commit carries a fixed, non-refundable `entry_fee` (fidelity-bond style), collected on both the crowd Commit and the ZK CommitDeposit, so filling an epoch with $k-1$ Sybils costs $(k-1) \times \text{fee}$ per epoch forever (the floor rolls forward until it is met). A `reward_bps` share accrues to an on-chain reward pool that funds a *dwell* reward: ClaimReward pays a drain-safe share proportional to how many distinct epochs a participant has committed into, so staying (which thickens future epochs) is what pays, not merely showing up. Dwell advances only through a real, fee-paying commit, at most once per epoch, so it cannot be minted for free. The ZK path deliberately ships no identity-linked claim; its anonymity-preserving equivalent (a dwell/age ZK proof) is designed, not implemented, and the docs say so.

6 Implementation and on-chain verification

Pure-Rust proving (no Node). The participant CLI (`mirror-cli`) proves in-process in pure Rust by default: `ark-circom` loads the same compiled circom `.wasm` witness calculator and the same committed

.zkey that snarkjs uses, computes the witness from typed Rust inputs, and **ark-groth16** over **ark-bn254** produces the proof using the snarkjs-compatible R1CS-to-QAP reduction, so the proof verifies under the same committed verifying key the on-chain program embeds. A legacy `--use-snarkjs` shell-out remains as an optional fallback. The CLI never self-submits: settlement is paid and signed by the rotating gasless relay, so the prove commands *emit* the settlement instruction bytes for the relay instead of sending them.

Trustless on-chain Groth16 (a check, not a committee). Both ZK paths verify their proof on-chain with **groth16-solana** (v0.2.0 byte layout) via the BN254 (**alt_bn128**) pairing syscalls, against a verifying key vendored into the program from **circuits/artifacts**. **proof_a** is emitted pre-negated so the on-chain path needs no ark serialization at runtime. This is the crucial contrast with a trusted-committee design: settlement does not trust an operator’s word that a member is valid; it checks a proof, and **verify()** (not **verify_unchecked()**) also rejects any public input not reduced modulo the field. The membership circuit (**circuits/membership.circom**, depth 20, **11522 R1CS constraints**, 4 public inputs, 41 private inputs) enforces leaf recomputation, Merkle inclusion, and the nullifier relation; public inputs are the fixed order [root, nullifierHash, actionHash, epoch]. Host and on-chain Poseidon (**light-poseidon** and **sol_poseidon**) are byte-identical and both match the circuit, pinned by a fixture cross-check test that reproduces the circuit’s **nullifierHash**, commitment leaf, and Merkle root exactly.

Compute cost. On-chain verification is cheap enough for per-membership settlement (about 170K to 500K CU for a 128-byte compressed proof); Section 8 gives the measured CU per flow on devnet.

Coordinator. **mirror-coordinator** maps its privacy jobs onto its modules: shared-epoch batching and the *k*-floor gate, honest **KAnon** accounting, the rotating gasless relay, and the normalized transaction profile (**cu_limit** = 400,000, priority fee = 10,000 micro-lamports). It is generic over a submitter seam, so the identical privacy-critical batching and floor logic runs against an in-memory submitter in tests and the real submitter in production, with no validator needed to test the decisions that matter.

7 Trusted setup: a verifiable multi-party ceremony

Groth16’s price for small, cheap on-chain verification is a per-circuit structured reference string whose sampling secrets (“toxic waste”) must be destroyed. We implement a distributable multi-party **phase-2** ceremony for both circuits (**crates/mirror-ceremony**, driven from **mirror-cli ceremony**), in pure Rust over the same arkworks stack the prover uses.

Phase 1 is imported, not generated. **ceremony start** takes a public perpetual powers-of-tau file and records its SHA-256, curve, power, and contribution count in the transcript; **ceremony inspect-ptau** parses the file’s own contributions section and prints every contributor’s self-declared name, so the list can be compared against the published record rather than asserted in prose. A file with fewer than two contributions is refused unless explicitly overridden. The initial phase-2 key is the output of **snarkjs groth16 setup**, which is deterministic: the same R1CS and the same powers-of-tau reproduce it byte for byte, so the start of the chain is re-derivable from public inputs.

A contribution moves only δ . With a fresh secret scalar s , δ_{G_1} and δ_{G_2} are multiplied by s and the δ -divided query vectors h and l by s^{-1} ; every other element of the key is unchanged. The final δ is the product of all contributions, giving the standard **1-of-N honest** property: the setup is safe if *at least one* contributor destroyed their scalar. Each contribution carries a Schnorr proof of knowledge of s with respect to the previous key’s δ_{G_1} , whose Fiat–Shamir challenge is bound to the running transcript hash, the contribution index, the contributor identifier, and the step’s kind and provenance; without extractable knowledge of the ratio, a contributor could publish a point derived from an earlier step and cancel an honest contributor’s randomness, and without the kind binding a closing beacon could be relabelled into an extra “contributor”. Entries are chained with SHA-256 over a canonical, length-prefixed serialization. A beacon is

final: the tool refuses to append after one, and verification rejects a chain that has a step after the beacon or more than one beacon.

Verification is reproducible by anyone. Given the transcript, the initial key, and the final key, `ceremony verify` recomputes the whole chain and checks the hash links, every entry hash, every proof of knowledge, a pairing same-ratio check $e(\delta_{G_1}^{prev}, \delta_{G_2}^{new}) = e(\delta_{G_1}^{new}, \delta_{G_2}^{prev})$ per step, exact recomputation of beacon steps from their published source, the endpoint digests, byte-equality of every δ -independent part of the key, and a batched pairing check that h and l were divided by exactly the accumulated ratio. It rejects a tampered delta, a forged or replayed proof of knowledge, a reordered chain, a truncated chain, a step appended after the closing beacon, and a beacon relabelled as a contributor; each rejection has a test, including the sophisticated variants where the adversary re-hashes the entire chain after tampering so only the algebra can object. Truncating *both* the transcript and the key is a legitimately shorter ceremony that the algebra cannot reject, which is precisely why the final transcript hash is the value a coordinator publishes.

Counting contributors honestly. The tool reports an *independent-contributor* count, not a contribution count. Contributions sharing an identity, a machine fingerprint, or a proof-of-knowledge nonce are merged; deterministic (fixed-seed) and beacon steps are never counted at all, because their scalars are reproducible by anyone. On a run where three self-asserted identities contributed from one machine, the tool reports **1 independent contributor from 4 steps**, not 4. We state its limits explicitly: identifiers and fingerprints are self-asserted, so this is a guard against accidental self-inflation and *not* a Sybil defence; it can also undercount when environments look alike, which is the safe direction. One limit is worth stating precisely, because it is information-theoretic rather than an implementation gap: a scalar carries no evidence of where it came from, so a verifier who does *not* hold the pre-committed beacon value cannot distinguish a public beacon scalar from a secret contributor's. Binding the kind into the proof of knowledge stops any third party from relabelling a published transcript, and passing the announced beacon value to `ceremony verify` turns the remaining case into a mechanical check; where that value is absent, the report says the check did not run rather than implying it did.

End-to-end. `ceremony prove-check` runs a ceremony, generates a real membership proof under the ceremony-produced proving key, and verifies it with the *exact* on-chain `groth16-solana` verifier against the ceremony-exported verifying key. In the exported key only `vk_delta_2` differs from the dev key, as phase-2 theory predicts, and a ceremony-key proof is correspondingly rejected by the old dev verifying key. `snarkjs groth16 verify` accepts the same proof under the exported `verification_key.json`.

What has not been done. No production ceremony has been *run*. The committed and deployed verifying keys still come from the insecure dev setup, and that remains true until a ceremony is run with external contributors and the program is redeployed with its key (Section 9).

8 Evaluation: public devnet deployment

Both soak suites were run against **public Solana devnet** (`api.devnet.solana.com`), a real, shared, public cluster. Every signature below is a real devnet transaction that resolves on the Solana Explorer; anyone can click through and verify it. This is devnet, not mainnet-beta: it proves the program deploys and every flow executes on a live public cluster exactly as it would on mainnet, without spending mainnet SOL. The SHA-256 of the on-chain program bytes equals the SHA-256 of the locally built `mirror_pool.so`, so the deployed program is exactly the source in this repository.

- **Program id:** `EezWdFrmHtR2PCuucUruvkgYB9HW3w2KskZNeYmXszBq`
explorer.solana.com/address/Eez...szBq (devnet)
- **Behavioral crowd settle** (4 identical actions, one atomic tx, one timestamp): `tx 4tcg...93NB`
- **Confidential private transfer** (`publicAmount == 0`, relay-only signer): `tx 2r6m...rtyv`

8.1 Behavioral soak: crowd + ZK settlement (17/17)

A fresh pool per run (`epoch_slots= 128`, `k_floor= 3`, `entry_fee= 106` lamports, `reward_bps= 2500`). Four participants commit the *same* `PlainTransfer` into one shared epoch, settled by one atomic gasless transaction (the sink is credited 4×10^7 lamports, i.e. $4 \times$ the 10^7 bucket); a ZK opt-in escrow of 5×10^7 is settled by an on-chain Groth16 `SettleZk`. The run makes one ZK deposit, so that settle exercises the mechanism (proof, binding, nullifier, escrow move) and not anonymity: its window holds a single deposit. The adversarial cases all fail closed on-chain: under-floor no-settle (`BelowKFloor`, custom error `0x2`), duplicate nullifier (`NullifierSpent`, `0x3`), re-settle (`EpochAlreadySettled`, `0x6`), mismatched recipient (`ActionHashMismatch`, `0xe`), and ZK replay (`NullifierSpent`). All 17/17 assertions passed.

flow	CU	flow	CU
<code>init_pool</code>	21,766	<code>underfloor_commit</code>	39,499
<code>crowd_commit</code> (x4)	39,647 to 41,006	<code>crowd_settle</code>	21,111
<code>zk_deposit_commit</code>	42,447	<code>zk_settle</code> (Groth16)	102,115

Table 6: Measured compute units per behavioral flow, finalized devnet transactions. The on-chain Groth16 membership verification (`zk_settle`) runs in about 102K CU, comfortably inside a transaction’s budget alongside the settlement work.

8.2 Confidential-value soak: shield / transfer / unshield (25/25)

Fresh pools per run: a main `ValuePool` and a fixed-denomination `ValuePool` (10^7 lamports). Amounts: shield 5×10^7 , hidden transfer 2×10^7 , withdraw 2×10^7 ; the relay fee (5000 lamports) is bound into ext-data. The assertions prove the headline claim on-chain: the vault is credited by the shield deposit and the value root advances with both output commitments inserted; the private transfer carries `publicAmount == 0` (verified as 32 zero bytes in the on-chain `Transact` data) and moves *no* public lamports, yet advances the root and spends its input notes; a recipient auto-discovers the payment note by scanning encrypted blobs; an unshield credits a fresh recipient and debits the vault by exactly the withdrawn amount; a wrong-denomination deposit is rejected on-chain (`DenominationMismatch`, `0x17`) and refused client-side; a mutated public input is rejected (`ProofVerificationFailed`, `0xc`); and conservation holds (vault = baseline + net deposit – net withdrawal). All 25/25 assertions passed.

flow	CU	flow	CU
<code>init_value_pool_main</code>	28,964	<code>transfer</code> (JoinSplit)	197,616
<code>init_value_pool_denom</code>	33,503	<code>unshield</code> (JoinSplit)	202,843
<code>shield</code> (JoinSplit)	196,837	<code>denom_shield_exact</code>	197,077

Table 7: Measured compute units per confidential flow, finalized devnet transactions. Each 2-in/2-out JoinSplit (7 public inputs) verifies on-chain in about 197K to 203K CU.

Test suite. Beyond the two live soaks, the repository ships **220 host tests** (`cargo test --workspace`, covering the Poseidon scheme cross-check against the circuit, the wire layout, `KAnon` accounting, the coordinator batching, *k*-floor and funding-round logic, ECIES notes, the trusted-setup ceremony, and the effective-*k* and funding-model functions; 7 further tests are environment-gated and skipped by default) and **44 on-chain program tests** (mollusk integration tests plus in-crate unit tests). `cargo build-sbf` is green and `cargo check --workspace` passes.

9 Honest limitations

A privacy tool that is vague about its non-goals is a trap. We list what still leaks and what we do not claim.

- **The DEPLOYED trusted setup is dev/test.** Both committed verifying keys (membership and the JoinSplit) came from a single-contributor, reproducible dev setup whose phase-2 entropy is a hard-coded public string, so their toxic waste is public. The keys verify and the program embeds them, but the system MUST NOT secure real value until a real ceremony output is exported and the program is redeployed with it. A multi-party phase-2 ceremony is implemented and runnable (Section 7); running one with external contributors is operational work that has not been done. This is a soundness caveat, disclosed here rather than buried.
- **Devnet, not mainnet.** The deployment proves execution on a live public cluster; it is not a mainnet-beta production deployment, and the system is not audited. Do not rely on it to protect real activity.
- **Anti-Sybil is economic, not a Sybil oracle.** Real k -anonymity assumes participants are economically distinct. The per-identity entry fee raises the price of flooding a round, but same-funding-source Sybils are not computable from the 32-byte commit stream and inflate `real_k`. The honest worst case: your anonymity is at least the number of participants you personally believe are independent, and at most `real_k`.
- **Effective- k is a model.** The numbers come from deterministic synthetic populations, a modeled Zipf common-funder distribution, and a modeled funding trace, not a live pool. A live-provenance loader is a named roadmap item.
- **The funding path is not wired and not soaked.** The unshield it is built on is soak-proven on devnet (25/25), but the layer above it is two disconnected halves: `fund-commit` proves an unshield to a fresh wallet and `prints` the emitted request, and `FundingRounds` is a unit-tested batcher type that would hold such requests to a round boundary. No shipped component ingests one into the other - the coordinator binary is an in-memory scheduler demo - so no funding round has ever run, and the two live soaks do not exercise this leg. Wiring the ingestion path and then soaking a funding round is named future work, not an implied result.
- **Funding provenance is reduced, not closed, and the headline number is conditional.** The funding mechanism turns a public edge into a matching problem, and the measurement says how much of that matching survives: 75.8% of nominal at $k=32$ under denominated batched rounds on the rules the protocol can actually enforce (a 24.2% residual), and 90.0% only if every participant also voluntarily dwells two rounds, which nothing in the protocol requires. Both figures assume 100% adoption; it drops to 42.4% at 50% adoption, and to 23.9% if participants pass value straight through a free-amount pool. The mechanism also does not hide that a wallet deposited into the value pool, and a participant whose deposit is itself the first hop out of a KYC'd exchange is still identified as a pool user.
- **The crowd path is behavioral, not cryptographic.** On the crowd path the participant signs their own action, so it stays attributable to that wallet, and the coordinator learns the participant-to-intent mapping while composing settlement (it reveals nothing the signed on-chain action did not). Use the ZK opt-in path for cryptographic who-initiated unlinkability. This is a disclosed distinction, not a hidden one.
- **Per-epoch, not cross-epoch.** An adversary with a prior over many epochs (intersection attacks) can narrow the set across epochs; neither path makes a cross-epoch claim, and mitigations are roadmap items.
- **Public boundary and membership.** Membership is public (like Tornado deposits); the confidential layer's shield-in and unshield-out magnitudes and the pool TVL stay public; a k -floor roll-forward is observable; and a participant who sweeps a fresh output into a clusterable wallet re-links themselves. Compliance and association-set tooling is future work.
- **The ZK path's set is per-window and per-amount, and no on-chain floor bounds it.** `SettleZk` publishes the epoch and the amount, so the covering set is that window's ZK deposits of the same amount, and the program will settle into a set of one: it reads neither `k_floor` nor any Epoch account. This asymmetry with `SettleEpoch` is deliberate. A crowd epoch below the floor *rolls forward* at no cost to anyone, whereas the ZK escrow's only exit is a `SettleZk` bound to one (recipient, amount, epoch), with no refund and no re-binding, so a settle-time floor would convert thin anonymity into permanently stranded escrow. Only the secret holder can produce the proof, so no third party can force a thin settle,

and the floor is enforced at proof time instead (the CLI refuses below `k_floor` unless explicitly waived). For the same reason recipient *freshness* is a client convention rather than a check: a dusted recipient address would otherwise strand the escrow. Both properties are pinned by mollusk tests.

- **The ZK escrow is a shared pot (v1 soundness).** The leaf is opaque, so nothing on-chain ties a settled amount to what its owner escrowed, and because both paths append the same leaf shape to one accumulator, a fee-only crowd `Commit` leaf satisfies the membership circuit too (a test demonstrates exactly this). A v1 pool must therefore not hold value it cannot afford to lose. Closing it needs a fixed denomination plus domain-separated leaves, a layout and circuit change; a partial fix would read like a full one.

10 Related work

mirror-pool is an independent clean-room implementation built from public techniques and papers. **Tornado Cash** established the commitment/nullifier/Merkle anonymity set for funds; mirror-pool reuses that accumulator shape but for *initiators* of behavior, and its behavioral core is the complement of a mixer (amounts public; on the crowd path the initiator *signal* is unattributable, on the ZK path the initiator itself is). **Tornado-Nova** contributed the 2-in/2-out Poseidon JoinSplit that the confidential-value layer implements clean-room. **Privacy Pools** added association-set membership to let honest users dissociate from illicit deposits; a mirror-pool association-set layer is named future work rather than a current claim. **Penumbra**’s batch swaps, where all swaps in a block clear at one price so per-user ordering is eliminated by construction, are the direct precedent for shared-epoch settlement: the per-actor signal removed structurally rather than through added noise [5]. Penumbra additionally shields the swap itself, which mirror-pool’s crowd path does not; that stronger property is what our ZK opt-in path buys separately. **Token-2022 confidential transfers** hide the transferred amount while leaving sender, receiver, and mint public, so the entity graph stays intact; they are the complement to the behavioral core (amounts hidden but behavior legible), and composing an amount-privacy primitive with a behavior-privacy one is exactly the “who + how much” point.

An honest tradeoff: transparent setups. mirror-pool uses Groth16, which needs a per-circuit trusted setup (Section 7; the deployed keys are still from the dev setup). Approaches built on STARKs or other transparent, post-quantum setups (for example Bulletproofs-style range proofs or STARK-based systems) avoid a trusted setup entirely, at the cost of larger proofs or higher verifier work. That is a real tradeoff, not a strict win for either side: Groth16 buys small, constant-size proofs and cheap on-chain pairing verification (the $\sim 100K$ to $200K$ CU measured in Section 8) in exchange for a ceremony; a transparent setup removes the ceremony risk in exchange for proof size or verification cost. mirror-pool’s choice is pragmatic for on-chain Solana verification today, and its trusted-setup limitation is disclosed rather than obscured.

11 Conclusion

mirror-pool composes two privacy axes that the literature usually treats separately: on-chain *behavior* and *how much* moved (a confidential-value JoinSplit). On the behavioral axis it ships two paths with two different strengths, and says which is which: synchronized-crowd batching removes the per-actor signal while leaving each action attributable to the wallet that signed it, and the ZK opt-in path, with an on-chain membership proof and no participant signature at settle, is the one that hides *who* initiated. It refuses to overstate anonymity, reporting a falsifiable metric backed by an adversarial harness and an information-theoretic effective- k model whose headline is given as the number the protocol can enforce rather than the number cooperative users can reach, and it verifies its zero-knowledge proofs trustlessly on-chain rather than trusting a committee. Every headline number in this paper is reproducible from the repository, and the whole system is deployed and browser-verifiable on public devnet. Appendix A makes each number checkable.

A Reproduce every number

All commands run from the repository root. The host workspace is one Cargo workspace; the on-chain program is a separate SBF crate excluded from it.

Build.

```
cargo build --workspace
cargo build-sbf --manifest-path programs/mirror-pool/Cargo.toml
```

Tests (220 host passing plus 7 environment-gated; 44 program).

```
cargo test --workspace
cargo test --manifest-path programs/mirror-pool/Cargo.toml
```

The 7 gated tests need build artifacts that are not committed (the compiled circuits, the powers-of-tau, the proving keys); they are `#[ignore]`d and additionally require `MIRROR_PROVE_LIVE=1`.

Trusted-setup ceremony (Section 7).

```
D=ceremony/membership
mirror-cli ceremony inspect-ptau --ptau circuits/pot16_final.ptau
snarkjs groth16 setup circuits/membership.r1cs \
circuits/pot16_final.ptau circuits/membership_0000.zkey
mirror-cli ceremony start --circuit membership --dir $D \
--r1cs circuits/membership.r1cs --ptau circuits/pot16_final.ptau \
--initial-zkey circuits/membership_0000.zkey
mirror-cli ceremony contribute --dir $D --id "alice@example.org"
mirror-cli ceremony contribute --dir $D --id "bob@example.net"
mirror-cli ceremony beacon --dir $D --id coordinator \
--source-hex <pre-committed value> --iterations-exp 16
mirror-cli ceremony verify --dir $D \
--r1cs circuits/membership.r1cs --initial-zkey circuits/membership_0000.zkey \
--beacon-source-hex <pre-committed value> --beacon-iterations-exp 16
mirror-cli ceremony prove-check --dir $D --out-dir /tmp/cproof
```

The last command proves the membership circuit under the ceremony-produced key and runs the on-chain `groth16-solana` verifier over the result; with `--out-dir` the same proof can be re-checked with `snarkjs groth16 verify`. The full guide, including what the independent-contributor count can and cannot detect, is in `docs/CEREMONY.md`.

Adversarial harness (the attack table and the effective- k table).

```
cargo run -p mirror-harness --release
```

Prints the Baseline-vs-MirrorPool advantage table (Section 5, FIFO $+77.54 \rightarrow -0.29$ pp at $k=16$) and the effective- k tables (funding-provenance alone: $k=32 \rightarrow 7.51$ for a public funding edge, 7.66 for pass-through shielded funding, 24.25 for denominated batched rounds at dwell 0 and 28.80 at the harness default dwell of 2; the funding-policy and adversary-strength ablations; the dwell and adoption sweeps; and the Sybil-dominance case collapsing to `real_k`). Deterministic (seed `0x4d4952524f52`), so any reviewer re-derives them bit-for-bit.

Public devnet soaks (the 17/17 and 25/25 assertion tables and the CU numbers).

```
PROG=EezWdFrmHtR2PCuucUruvkgYB9HW3w2KskZNeYmXszBq
cargo run -p mirror-soak --
--rpc-url https://api.devnet.solana.com --program-id $PROG
```

```
--epoch-slots 128 --k-floor 3 --zk-amount 50000000
cargo run -p mirror-soak --bin mirror-soak-value --
--rpc-url https://api.devnet.solana.com --program-id $PROG
--shield-amount 50000000 --transfer-amount 20000000 --denomination 10000000
```

The full reproduce recipe (fresh program keypair, single funded master payer, deploy, and re-run against a fresh devnet or a local mainnet-mirror validator) is in `docs/PROOF.md`, which also lists every finalized transaction signature.

Browser verification (no build required). Open the program address and the two representative transactions on the Solana Explorer:

- Program: <https://explorer.solana.com/address/EezWdFrmHtr2PCuucUruvkgYB9HW3w2KskZNeYmXszBq?cluster=devnet>
- Crowd settle: <https://explorer.solana.com/tx/4tcgNnhbZe7YKM26F6SnXW1WctASqVEqQ9W4v4NQ85fDfkYe3eq7YqCr4n7bEogm7U5By17Lkc5Tqa7jmZx693NB?cluster=devnet>
- Confidential transfer: <https://explorer.solana.com/tx/2r6mo2YrDtjJP8bnVqRJXaVDgVLkxaAmtKAHvgkYvWb8rchvd63QinFqbmuj5Gk2Se8M9g9ajcCkC5R6QNprtyv?cluster=devnet>

Where each claim lives. The threat model and the per-attack numbers are in `docs/THREAT_MODEL.md`; the two settlement paths and the account model in `docs/ARCHITECTURE.md`; the effective- k metric and its honest limits in `docs/EFFECTIVE_K.md`; the incentive economy in `docs/INCENTIVES.md`; the trusted-setup ceremony, how to contribute to one and how to verify somebody else's, in `docs/CEREMONY.md`; and the live results, signatures, and CU per flow in `docs/PROOF.md`.

References

- [1] *Deanonymizing Tornado Cash: an empirical analysis of mixed flows*. arXiv:2510.09433. FIFO temporal matching links up to 49% of deposits on small pools (+15 to 22pp over other heuristics); wallet fingerprinting via gas/priority-fee settings reduces the effective set by 37% alone; self-paid (non-relayer) gas is a dominant de-anonymization vector.
- [2] Y. Wang et al. *On How Zero-Knowledge Proof Blockchain Mixers Improve, and Worsen User Privacy*. arXiv:2201.09035. 27.34% (Ethereum) / 46.02% (BSC) anonymity-set reduction from composable heuristics; anonymity mining attracts privacy-indifferent users who degrade the set.
- [3] A. Serjantov and G. Danezis. *Towards an Information Theoretic Metric for Anonymity*. PET 2002. Defines the effective anonymity-set size as $2^{H(p)}$.
- [4] C. Diaz, S. Seys, J. Claessens, B. Preneel. *Towards Measuring Anonymity*. PET 2002. The companion entropy metric and the min-entropy / worst-case reading of anonymity as $1/\max_i p_i$.
- [5] Penumbra batch swaps. All swaps in a block execute at one clearing price, so per-user ordering and attribution are eliminated by construction and only the batch aggregate is revealed; the precedent for shared-epoch settlement.
- [6] Token-2022 Confidential Transfers (Solana Program Library). Hide the transferred amount while leaving sender, receiver, and mint public; the complement to mirror-pool's behavioral core.